



Autoregressive Image Modeling

OUESLATI Mohamed Amine
JEBARI Khaled

Summary

01

Introduction

02

MADE

03

PixelCNN

04

VAR

05

Results

06

Conclusion

01 Introduction

For an image represented as a vector

$$x = (x_1, x_2, \dots, x_D)$$

An autoregressive model writes the joint distribution as

$$p(x) = \prod_{i=1}^D p(x_i | x_1, \dots, x_{i-1})$$

01 Introduction

These models are usually trained via **maximum likelihood estimation**

$$\mathcal{L}(\theta) = -\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log p_{\theta}(\mathbf{x})] = -\mathbb{E}_{\mathbf{x}} \left[\sum_{i=1}^D \log p_{\theta}(x_i \mid x_{<i}) \right]$$

Sampling and generation:

1. Sample $x_1 \sim p(x_1)$,
2. Sample $x_2 \sim p(x_2 \mid x_1)$,
3. ... up to $x_D \sim p(x_D \mid x_{<D})$.

01 Introduction

We used **FID (Fréchet Inception Distance)** as an evaluation metric to compare images generated pixel-by-pixel by the model with real images, focusing on **distributional similarity** rather than likelihood

$$\text{FID} = \|\mu_r - \mu_g\|_2^2 + \text{Tr}\left(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2}\right)$$

02

MADE



02 MADE

MADE (Masked Autoencoder for Distribution Estimation)

Turn an autoencoder into a valid **autoregressive density model** by enforcing the factorization using **binary masks** on network connections.

Hidden layer

$$\mathbf{h}(\mathbf{x}) = g(\mathbf{b} + (W \odot M_W) \mathbf{x}) \quad (M_W)_{k,d} = \mathbf{1} [m(k) \geq d]$$

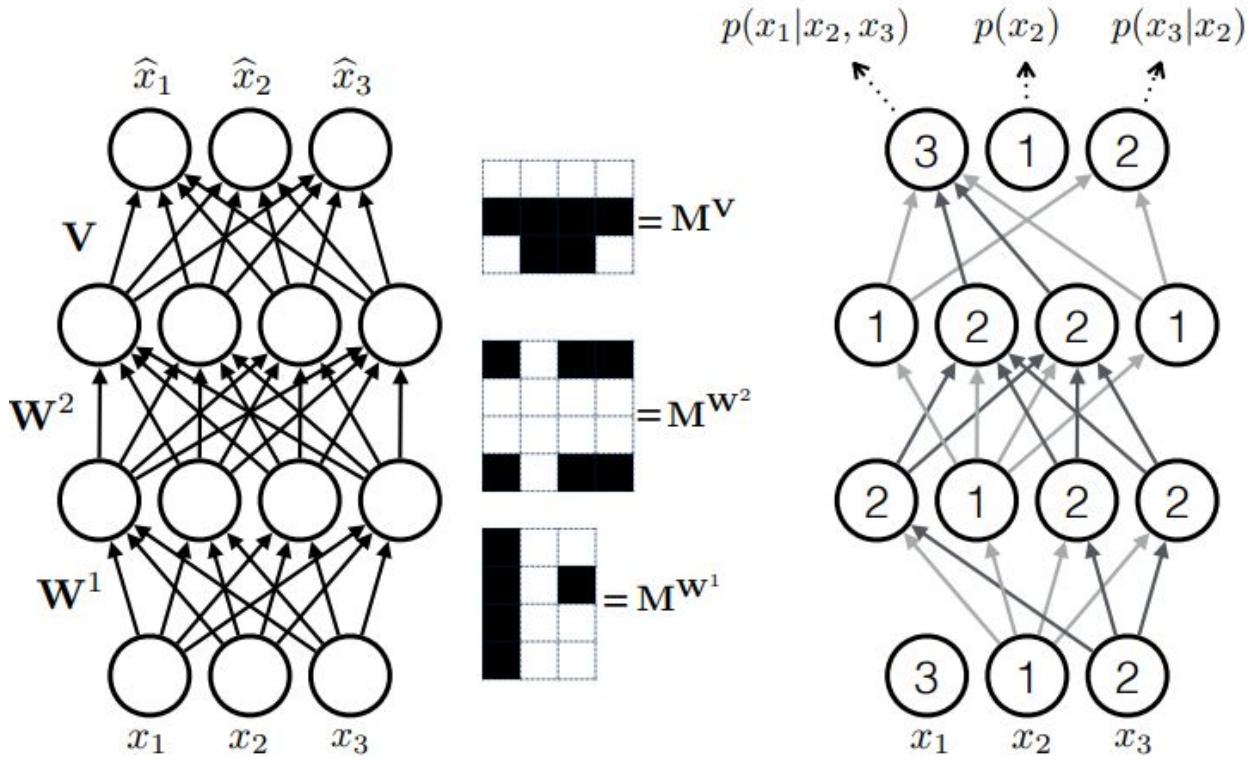
The mask removes illegal dependencies on future inputs

Output layer

$$\hat{\mathbf{x}} = \sigma(\mathbf{c} + (V \odot M_V) \mathbf{h}(\mathbf{x})) \quad (M_V)_{d,k} = \mathbf{1} [d > m(k)]$$

The mask ensures autoregressive structure

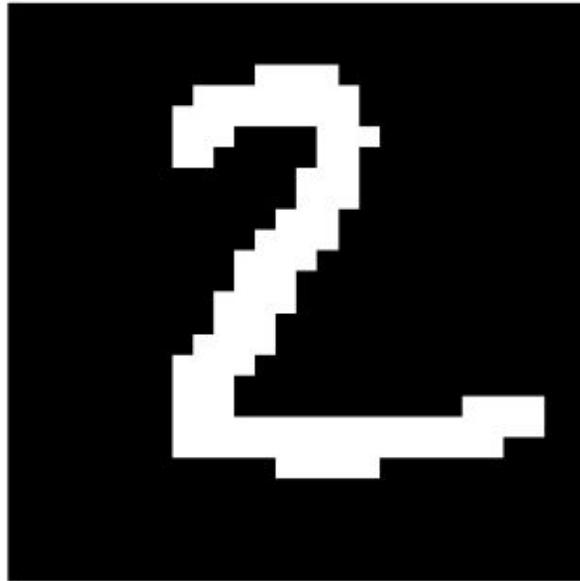
Each output $\hat{x}_d = p(x_d | x_{<d})$



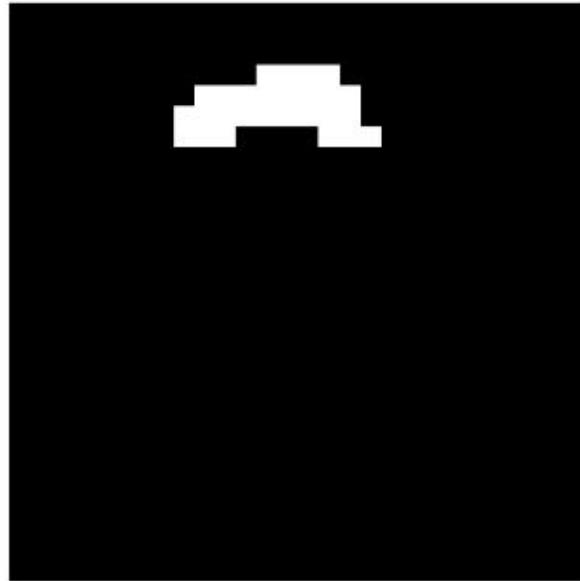
Autoencoder \times **Masks** \longrightarrow **MADE**

Binarized Mnist

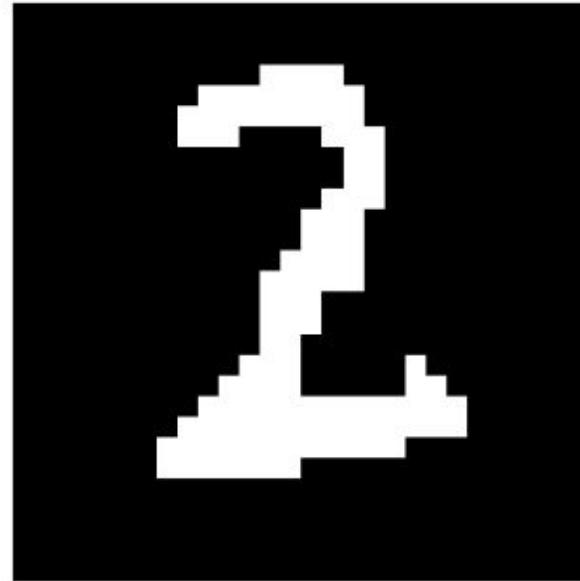
Original (Label: 2)



Observed (200/784 pixels)

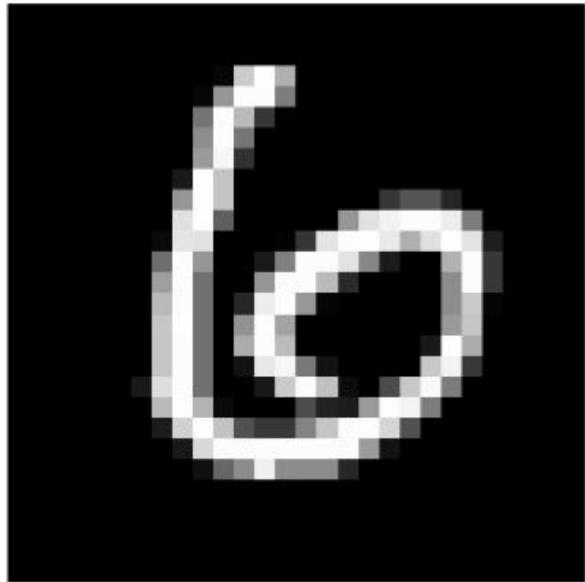


Completed by MADE

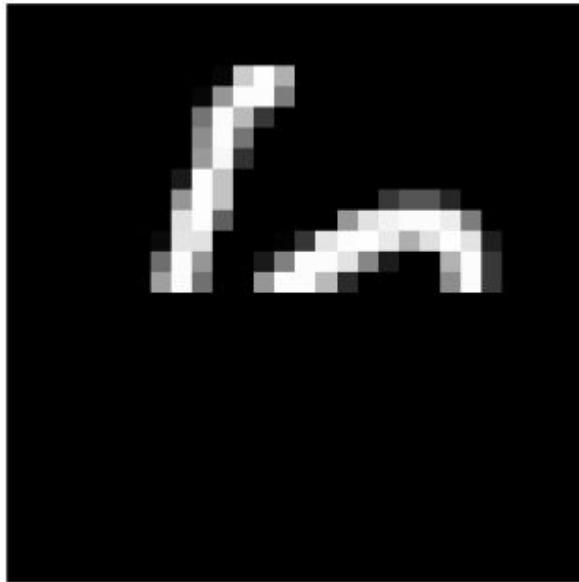


MNIST Grayscale

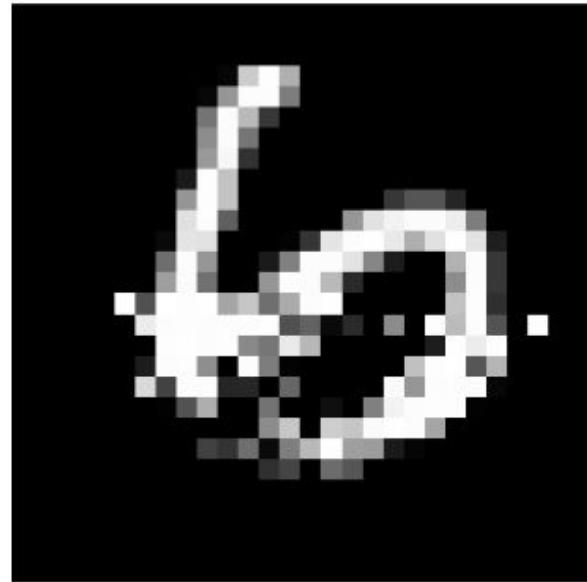
Original (Label: 6)



Observed (392/784 pixels)

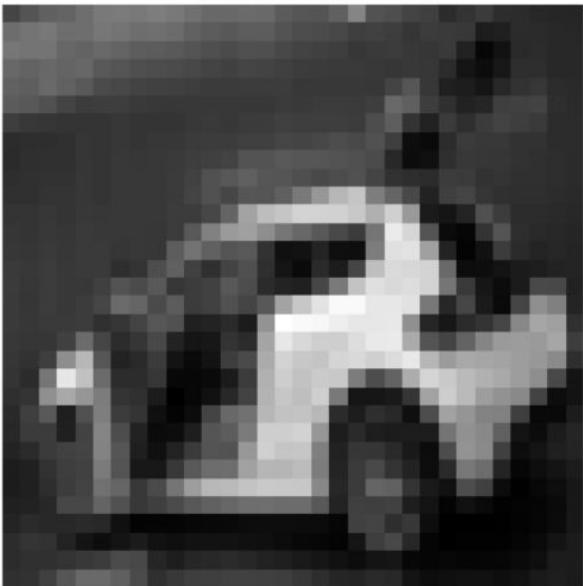


Completed by MADE

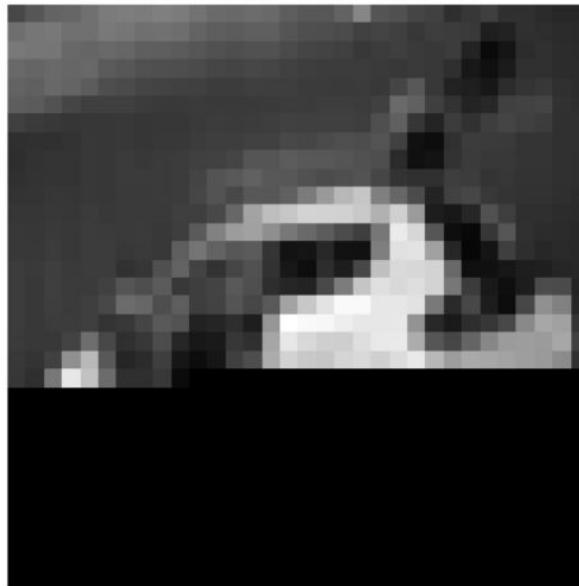


CIFAR10 Grayscale

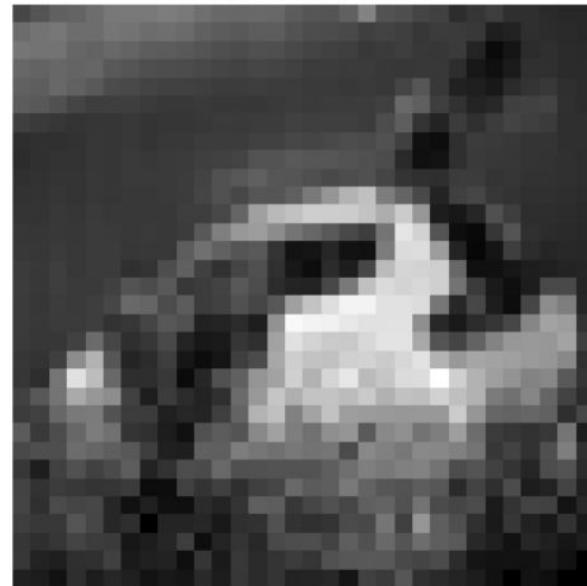
Original (automobile)



Observed (650/1024 pixels)



Completed by MADE



03

PixelCNN

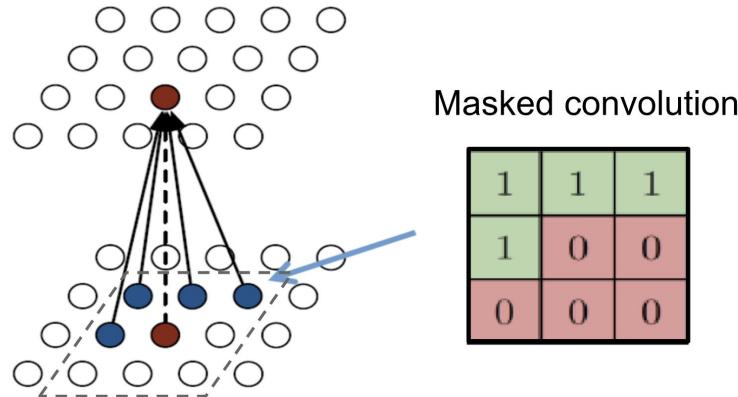


03 PixelCNN

PixelCNN uses **convolutional neural networks** with masked convolutions to ensure that the prediction of each pixel depends only on previous pixels

It defines a **binary mask** over the convolution kernel so that the receptive field of the convolution only includes “**past**” **pixels** (above or to the left)

Mask A prevents a pixel from seeing itself in the first layer, while Mask B allows self-conditioning through previous-layer activations in deeper layers



Masked convolution

1	1	1
1	1	1
1	0	0
0	0	0
0	0	0

1	1	1	1	1
1	1	1	1	1
1	1	0	0	0
0	0	0	0	0
0	0	0	0	0

(a) mask A

1	1	1	1	1
1	1	1	1	1
1	1	1	0	0
0	0	0	0	0
0	0	0	0	0

(b) mask B

MNIST

Autocompletion

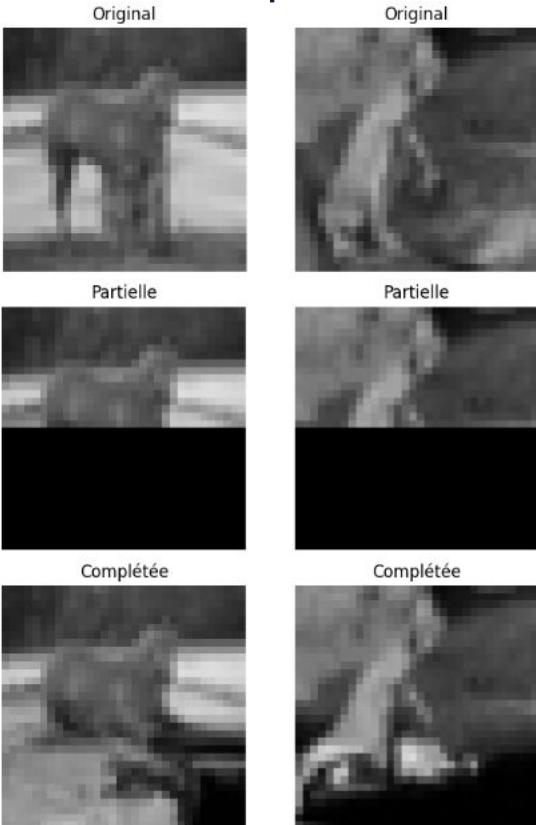


Sampling

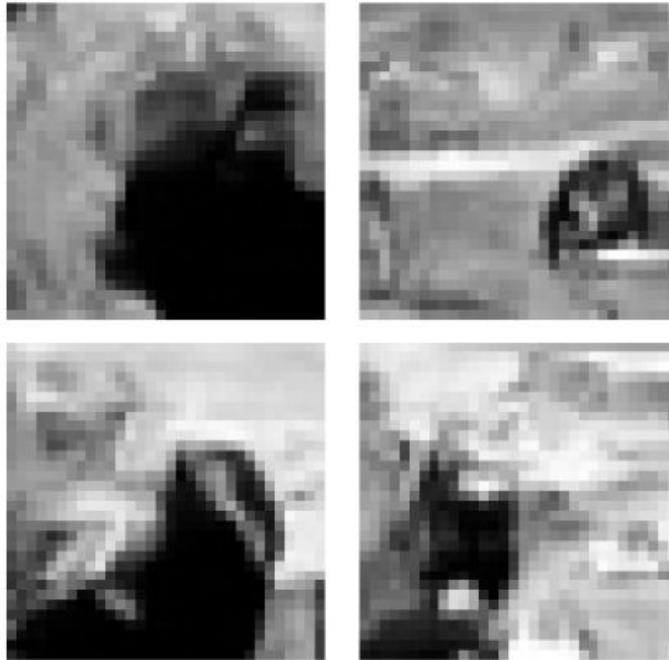


CIFAR10 Grayscale

Autocompletion



Sampling



04

VAR



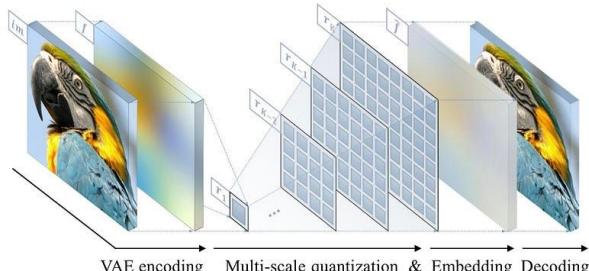
04 VAR

VAR (Visual Autoregressive Modeling) redefines the “generation order” by using a **coarse-to-fine, “next-scale”** strategy: generate low-resolution (coarse) token maps first, then progressively generate higher-resolution (finer) token maps, until reaching full resolution

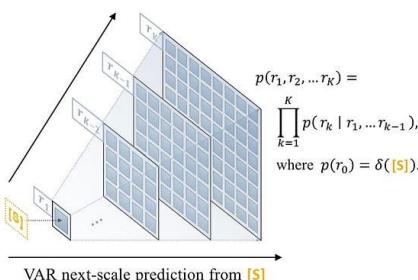
If we denote by T_1, T_2, \dots, T_S the token maps at increasing scales (from coarse to fine), then VAR defines the joint probability of the full image tokens as:

$$p(T_1, T_2, \dots, T_S) = p(T_1) \prod_{s=2}^S p(T_s | T_1, \dots, T_{s-1})$$

Stage 1: Training multi-scale VQVAE on images
(to provide the ground truth for Stage 2’s training)

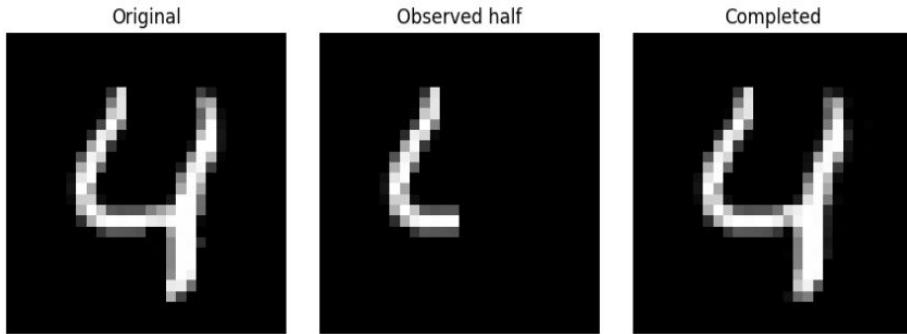


Stage 2: Training VAR transformer on tokens
($[S]$ means a start token w/ or w/o condition information)



MNIST

Autocompletion



Sampling



CIFAR10

Autocompletion



Sampling



05

Results



MADE

FID with respect to Number of Hidden Layers (here 500 neurons per layer)

Model / Hidden Layers	1	2	3
MNIST Binarized	53	58	64
MNIST Grayscale	270	274	313
Cifar10 Grayscale	258	262	271

Article loss results

MADE 1hl (1 mask)	88.40
MADE 2hl (1 mask)	89.59

Number of parameters ~ 0.8 M (CIFAR10)

MADE

FID with respect to Number of neurons (1 hidden layer here)

Model / nb neurons	500	1000	2000
MNIST Binarized	53	33	20
MNIST Grayscale	270	263	N/A
Cifar10 Grayscale	258	257	254

MADE

FID with respect to distribution

Model / nb neurons	500	1000
MNIST Grayscale Softmax	172	165
MNIST Grayscale Gaussian	270	263

PixelCNN

CIFAR10 Grayscale	170
CIFAR10 Grayscale with added layers	166
MNIST	27

Number of parameters ~ 1.5 M (CIFAR 10)

VAR - VQVAE

CIFAR10 Grayscale	194
MNIST	102

Number of parameters ~ 4M (CIFAR 10)

06

Conclusion

Bibliography

- [1] Mathieu_Germain et al. (2015) "MADE: Masked Autoencoder for Distribution Estimation" <https://arxiv.org/abs/1502.03509>
- [2] Aaron van den Oord et al.(2016) "Conditional Image Generation with PixelCNN Decoders" <https://arxiv.org/abs/1606.05328>
- [3] Keyu Tian et al.(2024) "Visual Autoregressive Modeling: Scalable Image Generation via Next-Scale Prediction" <https://arxiv.org/abs/2404.02905>

Generative AI Use

We used perplexity for bibliography search and to explain different parts in articles .
We used copilot with claude sonnet to implement Var model and to make some changes
in other models (like fid calculation)